Response to the Phase 1 feedback

Retrieval that Cannot See the Future

Pozorski, Jarosz, Florek, Cwalina

Remarks are quoted from the Phase 1 feedback (ellipses mark omissions). Section, table and figure numbers refer to `report.pdf` in this folder. "W" is the week of the plan in Sec. 4.

ID	Remark (quoted)	Status	Response	Where
F1	"Table 1 ... claims "every value names the dataset and the split". This is not true for several rows."	**Corrected**	Every row of Table 1 now names the split or says 'no split': Pikuliak 0.83 (all pairs, no split, same-language pool; the old label 'full pool' was also wrong), Ramponi (test partition of the 80-10-10 split), Vykopal (monolingual half, no split), Glockner (all claims, no split) and Ounis (49 judged topics, no split).	Table 1, Sec. 2
F2	"Schlichtkrull macro-F1 0.23/0.49 is labelled "AVeriTeC test". The paper's baseline analysis is on the development set."	**Corrected**	AVeriTeC numbers are development-set results; the test labels were not public. Label and text corrected.	Table 1, Sec. 2
F3	"Table 1 mixes scales: some values are 0–1 ... and others 0–100. Use one scale."	**Corrected**	One 0–1 scale; converted: Hardalov 90.3 to 0.903, Pikuliak 80.02/80.04 to 0.8002/0.8004, Vykopal 82.46 to 0.825, Liška 57/43% to 0.57/0.43, Glockner 69.7% to 0.697, Mu 91.5/47.2 to 0.915/0.472.	Table 1
F4	"Aschern's CheckThat! 2021 score as MAP@5 0.883 ... but as 88.2 ... Reconcile the two numbers or explain the difference."	**Corrected**	Both numbers are faithful to their sources: 0.883 is the CheckThat! 2021 overview, 0.882 is Aschern's value quoted in the table of Hardalov et al. Both are now stated, with the source of each.	Sec. 2, Table 1
F5	"The "No PFCR evaluation we found restricts the fact-check index per query" claim should be backed by a stated search protocol."	**Corrected**	Protocol written and run on 11 Nov 2026: OpenAlex, arXiv and forward citations of ten seeds, works since 2019; 1,738 records, 1,114 unique, 240 after keyword screening, 6 read at abstract level; none restricts the PFCR index per query. The closest (verification, case-law retrieval) are cited and have research notes.	Sec. 2; `research/protocols/causal-index-pfcr.yaml`
F6	"Class balance: "sparse positives (ranking)" is not a number."	**Corrected**	Measured on the raw files: gold FC per post and share of fact-checks with a linked post for v1, v2, SemEval and the Polish and English v2 slices.	Table 2

ID	Remark (quoted)	Status	Response	Where
F7	"Polish and English subset sizes are missing."	**Corrected**	Counts added from the raw files: for v2 and the SemEval test set, posts, dated posts, fact-checks, dated fact-checks and pairs per language (Polish v2: 1,450 posts, 1,406 dated, 1,550 pairs; English v2: 24,396 posts, 24,004 dated, 28,533 pairs; SemEval Polish 499 of 500 posts dated, 563 pairs, recounted from the raw file). For v1 the Polish and English counts are the published ones, because the v1 files we received have no language field and an empty text column and their ids differ from v2; we found no public statement of dated v1 posts per language.	Table 2, Sec. 3
F8	"The causal split ... and the 50 hand-coded failure annotations are deliverables. Each needs its own row with a TO COLLECT label."	**Corrected**	Rows added for the causal split (1,367 Polish and 22,600 English v2 queries, id lists kept confidential, code released) and for the 50 failure annotations: 25 Polish and 25 English queries, the seven failure codes, frozen before annotation, two annotators (Florek and Cwalina) with Cohen's κ, date 16 Dec.	Table 2, Sec. 4 (W5)
F9	"AMC-16K is RESTRICTED "to be requested", but no request date is given."	**Corrected**	The AMC-16K request was sent on 11 Nov 2026 and is entered in Table 2; the topic card's October request was the MultiClaim access, which is granted. Until AMC-16K arrives the judge validation (F25) cannot run, which Sec. 5 states.	Table 2, Sec. 4 (W1), Sec. 5
F10	"Licence column: "research use only" describes access terms. Name the actual terms document."	**Corrected**	Terms read from the Zenodo records (11 Nov 2026): v2 and SemEval limit access to researchers for non-commercial research (EU Directive 2019/790 research exemption); the v1 record gives no terms; none addresses id lists. We asked the dataset authors, who prefer to keep id lists confidential, so we will not publish them: deliverable (1) is the code that rebuilds the split from the restricted data. AMC-16K and MultiClaimNet name CC BY-NC-SA 4.0 and CC BY 4.0.	Table 2, Sec. 1 and 4 (deliverables), Sec. 5
F11	"No primary metric."	**Corrected**	Primary metric fixed before results: unconditional S@10, Polish–English difference-in-differences, random vs causal.	Sec. 4, Primary metric
F12	"The plan is not week by week. W1–W4 are blocks of about two weeks each."	**Corrected**	Weekly plan with one owner per task; only the remaining work is listed (the loader, pools and answerable rates exist).	Sec. 4
F13	"The fourth system is not scheduled ... "mmlw / LaBSE", which is undecided."	**Corrected**	mmlw-retrieval-roberta-large chosen and scheduled in W2.	Sec. 4
F14	"Figure 1 ... does not show the per-query date filter $P(q)$, the size-matched random control, or where the bootstrap enters."	**Corrected**	Figure 1 now shows $P(q)$ and the bootstrap; the size-matched control was already in it.	Fig. 1
F15	"RQ4 (LLM rerankers, "core+") has no metric, data or owner."	**Corrected**	RQ4 removed; stated as future work.	Sec. 1

ID	Remark (quoted)	Status	Response	Where
F16	"Undefined abbreviations: DPR, PoC, LDS, CLEF, TREC, KInIT, WUT, LaBSE, mmlw, GTR, RIET, rankSVM, "EN 2A"."	**Corrected**	Expanded at first use: DPR, PoC, LDS, CLEF, TREC, KInIT, LaBSE, rankSVM, GTR, QA, NDCG, EDA and GPU; 'EN 2A' is 'English Task 2A'. 'RIET Lab' is a team name; 'core+' no longer occurs.	Secs. 2–5, Table 4
F17	"(Ramponi et al., 2025) pages are 29045–29064."	**Corrected**	Ramponi pages corrected to 29045–29064.	Bibliography
F18	"(Chen et al., 2024) uses the arXiv title ... with the Findings of ACL venue."	**Corrected**	Chen et al. title corrected to the published one.	Bibliography
F19	"Citations used as sentence subjects are rendered in parentheses ... Use `\citet`."	**Corrected**	Citations as sentence subjects use `\newcite` (this template has no `\citet`).	Sec. 2
F20	"State explicitly that no per-outlet or per-author error ranking will be published."	**Corrected**	Sentence added: no per-outlet or per-author ranking.	Sec. 5
F21	"Restore the template credit line under the author block."	**Corrected**	Added under the title. It is in the course template; the first submission never had it, so it was added, not restored.	Title block
F22	"Define how tied timestamps, uncertain original publication dates, and missing earlier sibling fact-checks will be handled."	**Corrected**	Time is the earliest instance at full second precision; ties are excluded by $t_f < t_q$; margins of 1 and 7 days and conservative and liberal bounds are sensitivity runs; 100 URLs per language are audited; a post whose only earlier debunk is missing counts as unanswerable and is reported. Caveat: 138 of the 152 Polish answerable posts lie within a day of their gold fact-check.	Sec. 4, Dates and ties
F23	"Add the card's gap factors (time lag, lexical overlap, narrative drift) to the decomposition, or justify replacing them."	**Corrected**	Time lag, lexical overlap and narrative drift added to RQ2 as covariates of the decomposition.	Sec. 1
F24	"Name every tool and model in a corrected PDF ... give concrete tasks rather than role areas."	**Corrected**	The declaration names Claude Code with Claude Sonnet 5.5 and Claude Opus 5.5 and what they were used for; the contribution table lists named tasks per member from the repository history (branches main and feat/data-foundation), with hours for the Phase 1 report and this revision; the team confirmed the entries, including the text work that leaves no commits.	Sec. 7, Sec. 6, Table 4
F25	"No LLM judge appears anywhere else in the plan. Either remove it, or specify the judge model, its prompt and output labels, and the validation."	**Corrected**	Specified in Sec. 5: three open-weight judges disjoint from the evaluated systems, English prompt with the question and Yes/No answers, majority of the three, Cohen's κ per judge and for the majority on the 400 Polish AMC-16K pairs, accepted above 0.6 (threshold to confirm), and a fallback: sibling labels are reported unvalidated if AMC-16K has not arrived by week 4. Failure codes stay human.	Sec. 5, Table 2

ID	Remark (quoted)	Status	Response	Where
F26	"The strict causal pool will contain no gold item for most queries ... Report the answerable rate per language as the first Phase 2 result."	**Corrected**	Reported first: answerable posts per language and margin as a share of posts with a gold pair (v1 15.1%, v2 7.7%, Polish 11.1%, English 10.5%); 85–93% of gold pairs are post-first, so H1 holds.	Sec. 4, First result
F27	"The language-bias analysis ... is not cited. Cite it."	**Corrected**	Vykopal et al. (2026) cited to delimit the scope.	Sec. 1
F28	"Document restricted-data authorisations, and distinguish prespecified analysis from formally preregistered analysis."	**Corrected**	Authorisations: Paweł Florek requested access to the Zenodo records of MultiClaim v1 and v2, SemEval-2025 Task 7 and MultiClaimNet on 3 Nov 2026; approved on 4 Nov 2026. "Pre-registered" is replaced by "prespecified, not formally preregistered".	Sec. 1, Sec. 5
F29	"Test temporal, full-index and size-matched controls on equal query sets."	**Partly corrected**	Stated, and run for the strict and size-matched controls against the full pool on one query set per language for BM25 and multilingual-e5 (both languages) and mmlw (Polish) (Table 3). The bge reranker and matched historical leakage controls follow in W3.	Sec. 4, Table 3
F30	"Phase 2 requires a rebuttal ... and an extended SOTA table (criterion 2.5) that includes the works listed above."	**Corrected**	Table 1 now has the four suggested works that we cite and our proof-of-concept numbers; Urbano 2019 is cited for the permutation test.	Table 1, Sec. 4
Q1	"How do you handle tied publication dates, missing historical fact-checks and uncertain document timestamps?"	**Corrected**	Ties: excluded by strict $t_f < t_q$ at full second precision. Uncertain dates: earliest instance, margins of 1 and 7 days, conservative and liberal bounds, an audit of 100 URLs per language. Missing earlier siblings: the post counts as unanswerable and is reported. Caveat: most Polish answerable posts lie within a day of their gold fact-check.	Sec. 4, Dates and ties
Q2	"Can you show the effect of a strictly time-respecting index against both an unrestricted index and a size-matched control?"	**Corrected**	BM25 on one fixed query set per language against the full (unrestricted), size-matched and causal pools: Polish S@10 0.631, 0.484, 0.091; English 0.553, 0.398, 0.051. On answerable queries the causal pool costs 0.007 (Polish, 152 queries) and 0.027 (English), so nearly all of the drop is answerability; the Polish–English difference-in-differences of S@10 is -0.037 (95% interval $[-0.062, -0.010]$), and paired permutation tests give $p < 0.001$ unconditionally in both languages. Caveats: full-corpus document statistics, and 138 of the 152 Polish answerable posts lie within a day of the post, where time zones are unknown.	Sec. 4, Table 3
Suggested literature				
Vykopal 2026	"Cite it to delimit your contribution."	**Corrected**	Cited to delimit the scope (Sec. 1) and listed with its S@10 of 0.829 (Table 1).	Sec. 1, Table 1
Søgaard 2021	"random splits overestimate performance"	**Corrected**	Cited as the premise of the time-respecting split (Sec. 1) and listed in Table 1.	Sec. 1, Table 1

ID	Remark (quoted)	Status	Response	Where
Kazemi 2021, LongEval 2023	“missing from your multilingual PFCR section”; “a closer precedent for a time-respecting index”	**Corrected**	Both cited in the paragraph ‘Added after the Phase 1 review’ and listed with dataset and split in Table 1, last block.	Sec. 2, Table 1
Urbano 2019	“justifies the choice of paired bootstrap”	**Corrected**	The paper recommends the paired t-test or the permutation test and advises against the bootstrap-shift test; it fixes no resample count. We implemented a paired sign-flip permutation test (10,000 flips) for the p-values and keep the paired bootstrap over queries (1,000 resamples) for intervals only; both run in the proof of concept (code in the repository).	Sec. 4, Proof of concept, Table 3