Response to the Phase 1 feedback

Retrieval that Cannot See the Future

Pozorski, Jarosz, Florek, Cwalina

Remarks are quoted from the Phase 1 feedback (ellipses mark omissions). Section, table and figure numbers refer to `report.pdf` in this folder. "W" is the week of the plan in Sec. 4.

ID	Remark (quoted)	Status	Response	Where
F1	"Table 1 ... claims "every value names the dataset and the split". This is not true for several rows."	**Corrected**	Every row of Table 1 now names the split or says "no split". Pikuliak 0.83 is relabelled: all pairs, no held-out split, same-language pool (the old label "full pool" was also wrong). Glockner, AMC-16K and TREC have no split; the condition is named instead.	Table 1, Sec. 2
F2	"Schlichtkrull macro-F1 0.23/0.49 is labelled "AVeriTeC test". The paper's baseline analysis is on the development set."	**Corrected**	AVeriTeC numbers are development-set results; the test labels were not public. Label and text corrected.	Table 1, Sec. 2
F3	"Table 1 mixes scales: some values are 0–1 ... and others 0–100. Use one scale."	**Corrected**	All values on a 0–1 scale.	Table 1, Sec. 2
F4	"Aschern's CheckThat! 2021 score as MAP@5 0.883 ... but as 88.2 ... Reconcile the two numbers or explain the difference."	**Corrected**	Both numbers are faithful to their sources: 0.883 is the CheckThat! 2021 overview, 0.882 is Aschern's value quoted in the table of Hardalov et al. Both are now stated, with the source of each.	Sec. 2, Table 1
F5	"The "No PFCR evaluation we found restricts the fact-check index per query" claim should be backed by a stated search protocol."	**Corrected**	Protocol written (OpenAlex, arXiv, forward citations of ten seeds, since 2019, 11 Nov 2026) with counts at each stage; no PFCR work with a per-query pre-post index found. Two adjacent works (verification, case-law retrieval) are cited.	Sec. 2, `research/protocols/causal-index-pl`
F6	"Class balance: "sparse positives (ranking)" is not a number."	**Corrected**	Measured on the raw files: gold FC per post and share of fact-checks with a linked post for v1, v2, SemEval and the Polish and English v2 slices.	Table 2
F7	"Polish and English subset sizes are missing."	**Corrected**	Polish and English counts added: v1 from the paper, v2 and the SemEval test pool measured; dated posts and fact-checks stated; the Polish answerable count is 152 posts at margin 0.	Table 2

ID	Remark (quoted)	Status	Response	Where
F8	"The causal split ... and the 50 hand-coded failure annotations are deliverables. Each needs its own row with a TO COLLECT label."	**Corrected**	Rows for the causal split and the 50 failure annotations added: status TO COLLECT, size, label, annotators, date.	Table 2
F9	"AMC-16K is RESTRICTED "to be requested", but no request date is given."	**Partly corrected**	No request date was given and none was sent yet; the table now says so and W1 sends it. Note: the topic card's October deadline concerns the MultiClaim access, which is granted; the card sets no deadline for AMC-16K.	Table 2, Sec. 5
F10	"Licence column: "research use only" describes access terms. Name the actual terms document."	**Partly corrected**	Terms read from the Zenodo records (11 Nov 2026): v2 and SemEval limit access to researchers for non-commercial research (EU Directive 2019/790 research exemption), the v1 record gives no terms, and none of them addresses sharing id lists; we will ask the authors in writing before releasing the split, which is why the split row stays TO COLLECT. AMC-16K and MultiClaimNet name CC BY-NC-SA 4.0 and CC BY 4.0.	Table 2, Sec. 5
F11	"No primary metric."	**Corrected**	Primary metric fixed before results: unconditional S@10, Polish–English difference-in-differences, random vs causal.	Sec. 4, Primary metric
F12	"The plan is not week by week. W1–W4 are blocks of about two weeks each."	**Corrected**	Weekly plan, one owner per task. (W4 of the old plan was one week, not two.)	Sec. 4
F13	"The fourth system is not scheduled ... "mmlw / LaBSE", which is undecided."	**Corrected**	mmlw-retrieval-roberta-large chosen and scheduled in W2.	Sec. 4
F14	"Figure 1 ... does not show the per-query date filter $P(q)$, the size-matched random control, or where the bootstrap enters."	**Corrected**	Figure 1 now shows $P(q)$ and the bootstrap; the size-matched control was already in it.	Fig. 1
F15	"RQ4 (LLM rerankers, "core+") has no metric, data or owner."	**Corrected**	RQ4 removed; stated as future work.	Sec. 1
F16	"Undefined abbreviations: DPR, PoC, LDS, CLEF, TREC, KInIT, WUT, LaBSE, mmlw, GTR, RIET, rankSVM, "EN 2A"."	**Corrected**	Abbreviations expanded at first use. "RIET Lab" is a team name and kept; "core+" no longer occurs.	Secs. 2–5
F17	"(Ramponi et al., 2025) pages are 29045–29064."	**Corrected**	Ramponi pages corrected to 29045–29064.	Bibliography
F18	"(Chen et al., 2024) uses the arXiv title ... with the Findings of ACL venue."	**Corrected**	Chen et al. title corrected to the published one.	Bibliography

ID	Remark (quoted)	Status	Response	Where
F19	"Citations used as sentence subjects are rendered in parentheses ... Use `\citet`."	**Corrected**	Citations as sentence subjects use `\newcite` (this template has no `\citet`).	Sec. 2
F20	"State explicitly that no per-outlet or per-author error ranking will be published."	**Corrected**	Sentence added: no per-outlet or per-author ranking.	Sec. 5
F21	"Restore the template credit line under the author block."	**Corrected**	Credit line added under the title. (It was absent from the first submission, not removed.)	Title block
F22	"Define how tied timestamps, uncertain original publication dates, and missing earlier sibling fact-checks will be handled."	**Corrected**	Rule for several dated instances, ties (strict inequality), margins, bounds and missing sibling debunks added. Time is the earliest instance at full second precision, so ties are excluded by $t_f < t_q$.	Sec. 4, Dates and ties
F23	"Add the card's gap factors (time lag, lexical overlap, narrative drift) to the decomposition, or justify replacing them."	**Corrected**	Time lag, lexical overlap and narrative drift added to RQ2 as covariates of the decomposition.	Sec. 1
F24	"Name every tool and model in a corrected PDF ... give concrete tasks rather than role areas."	**Partly corrected**	Tools are named for Claude Code; the other members add theirs in the marked place (red TODO). The contribution table keeps role areas; concrete tasks are added in the Phase 2 version.	Sec. 7, Sec. 6
F25	"No LLM judge appears anywhere else in the plan. Either remove it, or specify the judge model, its prompt and output labels, and the validation."	**Corrected**	The judge is specified: three open-weight models, English prompt, Yes/No labels, validated on the 400 Polish AMC-16K pairs with Cohen's κ, threshold 0.6 (to confirm with the supervisor).	Sec. 5
F26	"The strict causal pool will contain no gold item for most queries ... Report the answerable rate per language as the first Phase 2 result."	**Corrected**	Reported first: answerable rates per language and margin (v1 15.1%, v2 7.7%, Polish 11.1%, English 10.5% of posts); 85–93% of gold pairs are post-first, so H1 holds.	Sec. 4, First result
F27	"The language-bias analysis ... is not cited. Cite it."	**Corrected**	Vykopal et al. (2026) cited to delimit the scope.	Sec. 1
F28	"Distinguish prespecified analysis from formally preregistered analysis."	**Corrected**	"Pre-registered" replaced by "prespecified, not formally preregistered". The access grants (Sec. 5) are those of the three Zenodo records.	Sec. 1, Sec. 5
F29	"Test temporal, full-index and size-matched controls on equal query sets."	**Corrected**	Stated, and run: the proof of concept scores one fixed query set per language under the random, size-matched and causal pools (BM25, MultiClaim v2; Table 3). Matched historical controls for leakage on the learned systems come with the other systems.	Sec. 4, Table 3
F30	"Phase 2 requires a rebuttal ... and an extended SOTA table (criterion 2.5) that includes the works listed above."	**Corrected**	Table 1 has the four suggested works we cite and our proof-of-concept numbers; Urbano 2019 is now a research note and is cited for the permutation test.	Table 1, Sec. 4

ID	Remark (quoted)	Status	Response	Where
Q1	"How do you handle tied publication dates, missing historical fact-checks and uncertain document timestamps?"	**Corrected**	See F22.	Sec. 4
Q2	"Can you show the effect of a strictly time-respecting index against both an unrestricted index and a size-matched control?"	**Corrected**	BM25 on one fixed query set per language: Polish S@10 0.631 (random), 0.484 (size-matched), 0.091 (causal); English 0.553, 0.398, 0.051. On answerable queries the causal pool costs 0.007 (Polish) and 0.027 (English); the unconditional drop is answerability. Caveat: full-corpus document statistics.	Sec. 4, Table 3
Suggested literature				
Vykopal 2026	"Cite it to delimit your contribution."	**Corrected**	Cited.	Sec. 1
Søgaard 2021	"random splits overestimate performance"	**Corrected**	Cited as the premise of the time-respecting split.	Sec. 1
Kazemi 2021, LongEval 2023	"missing from your multilingual PFCR section"; "a closer precedent for a time-respecting index"	**Corrected**	Cited and added to Table 1 with dataset and split.	Sec. 2, Table 1
Urbano 2019	"justifies the choice of paired bootstrap"	**Partly corrected**	The paper does not justify the bootstrap: it recommends the paired t-test or permutation test and discontinues the bootstrap-shift test. We cite it for a paired permutation test of the p-values and keep the bootstrap for intervals only.	Sec. 4